

Sardine: a Modular Python Live Coding Environment

Raphaël Forment
Université Jean Monnet (Saint Étienne, ECLLA)
raphael.forment@gmail.com

ABSTRACT

Sardine is a live coding environment and library for Python 3.10+ focusing on the modularity and extensibility of its base components (clocks, parser, *handlers*). Sardine has been designed to be easily integrated with existing *live-coding* environments as both a tool for experimentation and demonstration of various live coding techniques : temporal recursion, patterning, integration in various hardware and software setups. Although the tool is still in active early development, it has already been used in multiple public performances and algoraves. This paper is dedicated to the introduction of the **Sardine** system and the explanation of the main guidelines currently followed by contributors to the project. It will also present the preliminary results of our work through practical realizations that served as experimental validations during the early stages of development. Sardine already supports MIDI IN/Out, OSC IN/Out and *SuperCollider/SuperDirt* one-way communication through OSC.

1 Introduction

Sardine is a live coding library based on Python 3.10+ focusing on modularity and extensibility of its base components. Despite still being in early alpha stage, Sardine is extensively documented on a [dedicated website](#) providing installation guides, tutorials and media examples. Sardine is providing three main features linked together by the *FishBowl*, an environment handling synchronisation and communication between them:

- a *scheduling system* based on asynchronous and recursive function calls inspired by the concept of temporal recursion (Sorensen 2013). Calls can be scheduled in musical time either on an *InternalClock* or a *LinkClock* based on the *Link Protocol* (Goltz 2018).
- a *small and tidy number based pattern programming language* with support for basic generative and musical syntax (MIDI notes, polyphony, etc...), time-based patterns (*clock* and *absolute* time), handling of symbolic names.
- a *modular handlers* system allowing the creation and/or removal of various I/O (*OSC*, *MIDI*) or base components.

Sardine, by design, is in the direct lineage of previously released Python based libraries such as *FoxDot* (Kirkbride 2016), *Isobar* (Jones, n.d.) or the very recent *TidalVortex* (McLean et al. 2022). Initially conceived as a demonstration tool, Sardine partially emulates some selected features from the previously mentioned libraries or from the dominant live-coding *dialects* such as the *TidalCycles* (McLean 2014) rhythmical mininotation or the *Sonic Pi* (Aaron 2016) imperative writing syntax. Sardine is designed as a general *agnostic* framework for approaching live coding using Python. Thus, the library is aiming to support different writing paradigms and different approaches to live performance based on the manipulation of source code. The reliance on regular Python asynchronous functions for scheduling and music writing has for consequence that Sardine is particularly suited to let each developer-musician follow its own personal coding style, ensuring a blank slate for experimentation. Furthermore, Sardine design has been strongly influenced by Andrew Mc Pherson's and Koray Tahiroğlu concerns about the *idiomatic patterns* (McPherson and Tahiroğlu 2020) of usage enforced by computer music softwares, pushing users to repeat and strictly follow preferred patterns of usage.

The version hereby presented, labelled as 0.2.0, is offering a first-look into the complete intended design for the library. It features a near complete rewrite over the 0.1.0 version previously used by members of the french live coding scene and by the first global Sardine users. It features two different clock implementations, multiple handlers for I/O (*MIDI*, *OSC*, *SuperDirt*), a robust asynchronous temporal recursive scheduling system and a reimagining of the *'Player'* system previously introduced by *FoxDot*(Kirkbride 2016). Sardine originality lies in its temporal model, strongly anchored in Python's default mechanisms for asynchronous programming. The use of the *asyncio* library is offering a variant to other threaded musical clocks featured by past Python based live coding libraries. Sardine also presents itself as an



Figure 1: *Sardine first algorave in Lorient (France), 2022, October 13th. Photography: Guillaume Kerjean.*

agnostic and minimal tool modular enough to be integrated into any live-coder tooling and setup, capable of handling any Python-based scheduling duties. As such, Sardine has been already successfully integrated in various laptop-based performance setups involving audio, video and hardware components.

On the technical side, Sardine has been developed entirely using the Python programming language, with few libraries depending on C++ code through bindings to external libraries. Despite the known short-comings of Python for interpreted conversational real time programming (incomplete support of dynamic programming, slowness relative to other interpreted languages), we do believe that this language is suitable for the implementation of a live coding library. The large collection of available libraries and modules and the popularity of the language ensures the affordance of good tooling and rich customization and Sardine integration options into different text editors, running environments, etc... Sardine already takes advantage of a thorough ecosystem of libraries focused on data *input/output*, network communication and text manipulation. Moreover, thanks to its lightweight and clear syntax, Python can be read by programmers coming from different domains with a minimal adaptation time, making it a convenient platform for collaboration and experimentation over the implementation of bespoke features needed by performers.

In the present article, we will introduce the *Sardine* system by detailing its goals (1) and base implementation (2). By doing so, we hope to highlight the basic principles of its inner working while providing some context on the current direction taken by the project and by its users (3).

2 Methodology and objectives: a framework for exploring live-coding in Python

Sardine is born out of a curiosity for the implementation of similarly featured Python-based live-coding libraries such as [FoxDot](#), [Isobar](#) or the very recent [TidalVortex](#) (McLean et al. 2022). At its inception, the Sardine project was thought as an attempt to provide a functional but barebones live coding library for demonstration purposes in a dissertation manuscript; a library capable enough for showing the impact of design and implementation choices on the possibilities of musical expression and on the expressiveness offered by a live coding environment. Therefore, a particular attention has been given to reproducing or *at least* paving the way for the reproduction of different coding styles and representation of timed musical information. The base design has quickly evolved after the first initial public tests. It has been decided to aim for an increased modularity of the system in order to support and maximise the *input* and *output* options offered by Sardine. This has allowed for the quick integration of the tool with other neighbor interfaces and live coding environments.

The development of Sardine began initially in a period of frantic collaborations and joint performances with the parisian *Cookie Collective* (Collective 2016) and the Digital Audio Community from Lyon (*th4*, *ralt144MI*, etc...). Stemming from the *demoscene* and shader-coding scene, the *Cookie* is known for its complex multimedia performances, each member relying on bespoke hybrid audio-visual setups ranging from low end computing devices to complex synthesizers and circuit-bended video mixers. It is also known for working in an improvised manner, customising its setup for each venue depending on the audience needs and expectations. The need to adapt and customize the live coding interfaces already in use to the needs of each performance and each artist gave rise to the idea of creating a modular interface that

could be used and mastered by all the members of the collective, while allowing for jam-ready synchronisation with other musicians and live-coders. The splitting of Foxdot’s development into several competing branches reinforced the need for a customizable and easily editable Python interface by the community. Due to the open-ended nature of the development process, Sardine has been gradually shifting towards its current modular architecture, allowing each performer to refine the nature of the *inputs* and *outputs* controllable by the system, from simple MIDI note output to more convoluted custom Sysex and OSC message support. The invaluable help and expertise from John Phan has allowed for a complete deep rewrite of every base mechanism. The completion of the program rewrite marks the beginning of a new stage in the development process, focused on introducing new features (new audio backends? better OSC support?) and improving existing ones (parser?). This process is managed in an *ad-hoc* manner, by encouraging users to propose ideas and contribute to an extensively documented codebase.

3 Sardine implementation




Figure 1: Software layers of the Sardine system stack.

Sardine is implemented and distributed as two complementary Python modules: *sardine* and *fishery*. They work hand in hand in a layered and coordinated fashion. *fishery* provides amendments to the default Python asynchronous REPL¹ and constitutes the entry point for the Sardine system, accessible by typing `python -m fishery` or simply `fishery` right after install. *fishery* is nothing more but a slightly modified version of the base Python asynchronous REPL. Importing it also imports *sardine* and will *de facto* starting a new playing session. As an helper for new users, a terminal based configuration client (*sardine-config*) is also provided and can be used to setup various options before starting *fishery*. Configuration files are stored in a default standard location depending on the OS currently in use (e.g. `.local/share` on UNIX systems). Configuration files include a general JSON configuration file, a `.py` file usable to load user-specific Python code at the start of each session and files needed to properly configure a *SuperCollider*/*SuperDirt* session. This architecture, despite its initial complexity, is being used to make Sardine more accessible to novice users who may not be familiar with using the command line and Python development tools. Files relative to Sardine will be kept in a single configuration folder. Note that modularity of the system is greatly encouraged, with many *input* and *output* components being disabled by default. This makes the installation of any audio-backend like *SuperCollider* entirely optional, being considered more as a target than a dependency.

Thanks to the generally great IDE support for Python, *Sardine* is not shipping with its own text-editor or dedicated text editor plugin. *Sardine* has been tested with third-party code writing tools such as *Atom*, *VSCode*, *Emacs*, *Vim/Neovim* or even *Jupyter Notebooks*. Each one of these text editors generally support the spawning of an asynchronous REPL and the piping of code from a text buffer to a running interpreter. The setup process for each one of these interfaces generally relies on the installation of a simple general-purpose Python plugin². This state of fact has lead us to consider the Python interpreter as a code receiver and monitoring tool mainly used to mirror useful information to the user, such as the state of the *SuperCollider* sub-process, of the event loop and *runners*, etc... Every other operation is directly handled by calls internal to a *Sardine* session.

Reliance on any audio backend can / will require the boot of another application. For the time being, only *SuperCollider* and *SuperDirt* are natively supported by their own *Sardine* components. Even though the installation of these backends is still necessary for users willing to use them, integration is done in such a way that there is no need – later on – to actively take care and monitor manually any of these dependencies. A basic API to *SuperCollider* and *SuperDirt* is offered through the `SC.send()` function, allowing to run arbitrary `scLang` code in the subprocess session. The addition of more automatically-managed audio-backend *subprocesses* is actively planned and will be explored in the coming months (deeper *SuperCollider* integration, *CSound* backend, etc...). Clever combination of *Sardine* provided functions is already allowing some amount of customization for patterning hardware and software synthesizers through MIDI or OSC.

¹*Read, Eval, Print, Loop*: mechanism used by most interpreted languages to quickly process user input from the command line.

²The process for setting up various interfaces is extensively detailed on [Sardine Website](#)

Being packaged as a regular Python module, Sardine makes use of the `pyproject.toml` module configuration and packaging format defined by PEP 660. This has for advantages that no third party tool is currently required to install Sardine other than a base *complete* installation of a modern (3.10+) Python runtime. However, one must note that the package is not, at the time of writing, fully installable in the binary ‘*wheels*’ format generally favored by Python developers and users alike. This has to do with the problematic packaging of some C++ external dependencies used by Sardine to process various I/O processes. Future versions will hopefully be tightly packaged and served through the central *Pipy* package distribution system. Most users will still have to install part of the compilation toolchain (CMake and any compatible C++ compiler) to be able to manually compile these dependencies for the target system.

3.1 Event loop and scheduling System

3.1.1 Event loop

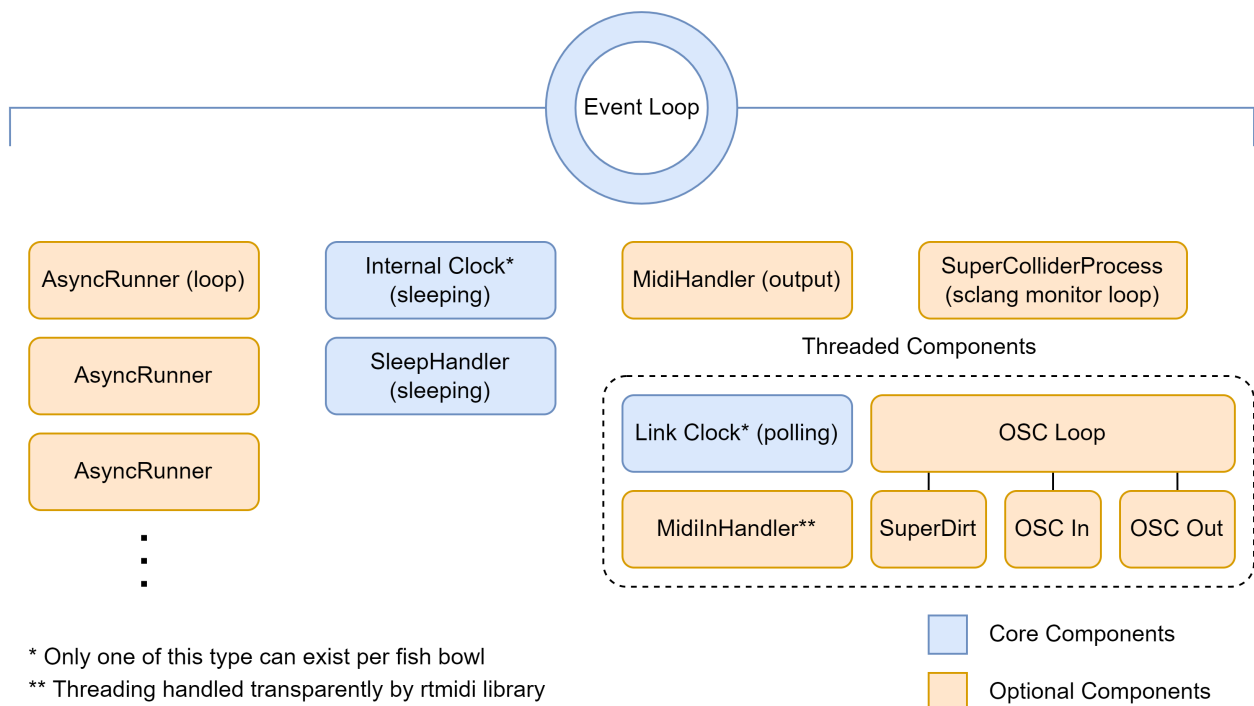


Figure 2: Architecture diagram of the customised asynchronous event loop.

Sardine is making use of the asynchronous programming features offered by Python. More specifically, Sardine takes advantage of the not well known `asyncio` REPL prototype introduced by Python 3.8 (Selivanov, n.d.). The `UVLoop` (Stack 2016) drop-in replacement event loop is also being used in order to speed up the scheduling of asynchronous calls. Several hot-patches to the asynchronous loop have been introduced by John Phan (*thegamecracks*) in order to make its behaviour consistent on every major OS platform. Sardine is laid out as a series of abstractions built on the base loop, making it aware of tempo and timing. Sardine *clock* (either the internal or link) clock automatically starts whenever the system is imported but pure asynchronous calls can still be handled even if the clock is being stopped.

The consistency of the asynchronous clocks is being covered by tests (in the `tests/` folder) and has been checked to be *on-par* with the alternatives offered by similar more widely used threaded clocks. Development of such a feature has proven to be a difficult technical challenge due to the specificity of the task and of the relatively obscure inner workings of internal OS’s schedulers. Threaded components are still used for various *I/O* operations in order to lighten the load of the event loop and to alleviate the temporal cost of message processing. Note that many *Sardine* components are entirely optional and can be activated on demand by the user. Only the `clock`, `AsyncRunners` and `SleepHandler` constitute the core abstractions needed over Python `asyncio` loop. Basing the custom event loop on top of the Python asynchronous interpreter is allowing for the evaluation of any top-level asynchronous `await` instructions that would be forbidden by the main interpreter. It must be noted that Python `asyncio` features have their own logic for every major OS and that some differences can be noted when testing under different systems.

3.1.2 Scheduling

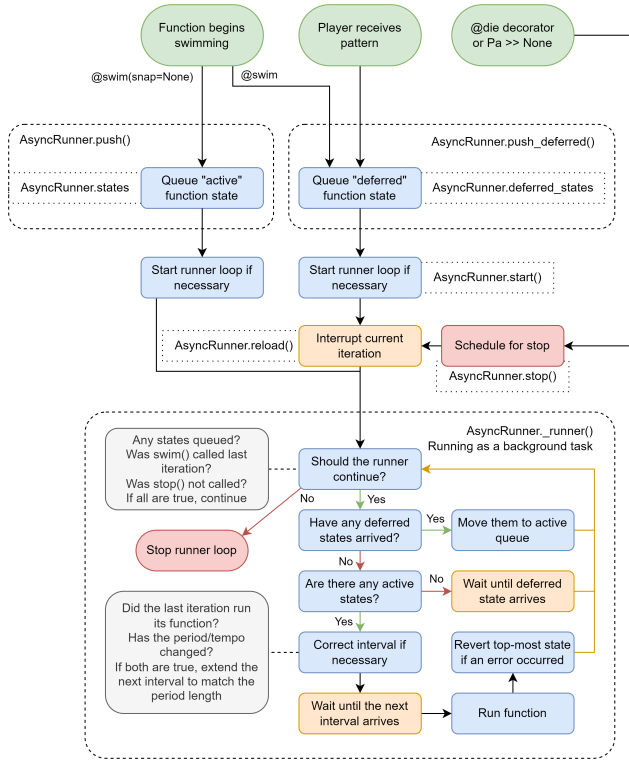


Figure 3: Lifetime of an asynchronous 'swimming function'.

Python is known to be a language that doesn't have native support for tail-call recursion (Rossum 2009b, 2009a), making the infinite recursion of a function a delicate task. To properly support this central feature, a complex system based on John Phan's AsyncRunners has been developed and is used as the basis for every repetitive operation (such as a pattern) scheduled with Sardine. In the spirit of the metaphor followed by the whole program, a temporal recursive function is called by the development team a *swimming function*, and is labelled in code as an AsyncRunner. A *swimming function* can be started using the @swim decorator³, stopped using the @die decorator and can receive updates all along its lifetime on the scheduler.

Decorating a Python function is enough to push a given synchronous or asynchronous function to the scheduler, making it repeat every p (for period), a time measured in beats relative to the clock currently in use. The content of a given function will be re-evaluated for every recursion cycle and state can be preserved either by passing arguments to a subsequent call or by relying on global state. *Swimming functions* are a powerful construct for building abstractions dealing with time, code re-evaluation and dynamic lifetime management of code components. Iterators, for example, can be built by incrementing a variable passed as argument. Random generators can be built by calling a simple native random function whose result will be dynamically updated for each recursion.

Swimming functions will automatically start *on-the-beat*. The start of a function can target a specific point in musical time by specifying a special snap argument that is understood as an offset, in beats, from the beginning of the next bar. The period argument of a given function is the only required argument for a function to be considered as a valid *swimming function*. Every other component of the Sardine system works on the assumption that its evaluation context will be the *swimming function*. They can receive any arbitrary Python code and/or call the various players defined by the Sardine system to properly handle I/O operations. Thus, the prototype of a basic musical function using the base model looks like:

```

@swim # swimming decorator (swim or die)
def swimming_function(p=0.5, i=0): # p: (period), i (custom iterator)
    print('I am swimming in time.')
    D('bd, hh, cp, hh', i=i) # call to the 'Dirt' SuperDirt interface.
    ... # user specified code
    again(swimming_function, p=0.5, i=i+1) # recursion callback with argument passing
    # remove the call to again() to stop the recursion from happening, stopping the runner.

```

Multiple abstractions can be built on top of the basic *swimming function* mechanism, allowing for a terser user-facing syntax. We believe that building abstraction on top of the *swimming function* is helpful to allow newcomers to get a grasp on the temporal model offered by the system. The FoxDot's inspired *surfboard* mechanism is currently the only abstraction available demonstrating this principle. It automatically handles its own scheduling logic and also provides iterators needed by the *Senders* that we will detail later on:

```

Pa >> d('bd, hh, cp, hh', p=0.5) # Terser version of the above swimming function

```

³Decorators in Python are used to add a behaviour to an object without modifying the base object itself.



Figure 2: On the left pane, Sardine being imported through fishery*. On the right, the sardine-config configuration client.

3.1.3 Central dispatch environment

Every component of the system can talk or access the data held by any other component through the central 'FishBowl' mechanism. The environment has been implemented as a message dispatcher allowing each component to subscribe to it through *hooks* and making each component capable of answering to dispatch messages (such as `bowl.dispatch('stop')` for stopping the clock). Some components are to be considered as hard dependencies, such as the clock and the parser as they provide the basic mechanisms needed by every other modular component to properly function. Every other can just take advantage of the central environment to provide additional features or logic (such as a midi output, an osc input, or any other user-defined construct).

3.2 Sardine Pattern Language

A small patterning language has been developed for Sardine using the [Lark](#) parsing toolkit. Defined as a LALR parser, the syntax of the language is best described as a list-based calculator capable of dealing with basic MIDI note definition, custom chance operators and other composition tools.

3.3 Players and Handlers

Description of the event based system. How to define an handler, what is an handler, etc...

Demo of the SuperDirt handler, etc...

4 Sardine usage

Basic facts about the usage of Sardine in various text editing environments + how to install and handle a Sardine installation.

4.1 Algorave and performance

Zorba, Lorient, example code taken from performances.

4.2 Controlling Legacy MIDI Synthesizers

Rémi Georges usage of Sardine: controlling legacy synthesizers along with TidalCycles, etc...

4.3 Usage of Sardine at the II Laboratory

Projects involving the Magnetic Resonator Piano, Boids, etc...

5 Project directions

5.1 Packaging and distribution

Distribution and release for Python 3.11 with updated C++ dependencies whenever possible. Distribution on Pypi when it'll be bug free, etc...

5.2 Opening up for collaboration

Documenting, section about the website and integration of the Sardinopedia.

5.3 Creation and performance

6 Conclusion

Call for contributors, etc...

7 Acknowledgments

I warmly thank my thesis supervisors Laurent Pottier and Alain Bonardi for their support and advice in the creation of this tool. I thank the doctoral school *3LA* from the University of Lyon for the funding it provided to this research. I extend my thanks to the musicians and friends who allowed me to take *Sardine* on stage and to present it to a wider audience these few last months: the *Cookie Collective*, Rémi Georges, etc...

References

- 10 Aaron, Sam. 2016. "Sonic Pi—Performance in Education, Technology and Art." *International Journal of Performance Arts and Digital Media* 12 (2): 171–78.
- Collective, Cookie. 2016. "Cookie Collective." <https://cookie.paris/all/>.
- Goltz, Florian. 2018. "Ableton Link—a Technology to Synchronize Music Software." In *Proceedings of the Linux Audio Conference*, 39–42.
- Jones, Daniel John. n.d. "Isobar." <https://github.com/ideoforms/isobar>.
- Kirkbride, Ryan. 2016. "Foxdot: Live Coding with Python and Supercollider." In *Proceedings of the International Conference on Live Interfaces*, 194–98.
- McLean, Alex. 2014. "Making Programming Languages to Dance to: Live Coding with Tidal." In *Proceedings of the 2nd ACM SIGPLAN International Workshop on Functional Art, Music, Modeling & Design*, 63–70.
- McLean, Alex, Damian Silvani, Raphaël Forment, and Sylvain Le Beux. 2022. *TidalVortex Zero*. Zenodo. <https://doi.org/10.5281/zenodo.6456380>.
- McPherson, Andrew, and Koray Tahiroğlu. 2020. "Idiomatic Patterns and Aesthetic Influence in Computer Music Languages." *Organised Sound* 25 (1): 53–63.
- Rossum, Guido Van. 2009a. "Final Word on Tail Calls." Blog.

———. 2009b. “Tail Recursion Elimination.” Blog.

Selivanov, Yuri. n.d. “Implement Asyncio REPL.” <https://github.com/python/cpython/issues/81209>.

Sorensen, Andrew. 2013. “The Many Faces of a Temporal Recursion.” http://extempore.moso.com.au/temporal_recursion.html.

Stack), Yury Selivanov (Magic. 2016. “UVLoop.” *GitHub Repository*. <https://github.com/MagicStack/uvloop>; GitHub.